

# General SQL Topics

## 1. INTRODUCTION TO SQL

- **SQL:** Structured Query Language, used to interact with relational databases.
  - **Relational Database:** Stores data in tables (rows & columns).
  - **Common RDBMS:** MySQL, PostgreSQL, SQLite, Oracle, SQL Server.
- 

## 2. SQL DATA TYPES

- **INT:** Integer numbers
  - **VARCHAR(n):** Variable-length string with a maximum of n characters
  - **CHAR(n):** Fixed-length string
  - **DATE / DATETIME:** Date or timestamp
  - **DECIMAL(p, s):** Exact numeric value with precision and scale
  - **BOOLEAN:** True or False
  - **TEXT:** Long-form string data
- 

## 3. BASIC SQL COMMANDS

### ► Create Table

```
CREATE TABLE Employees (  
  id INT PRIMARY KEY,  
  name VARCHAR(100),  
  salary INT,  
  department VARCHAR(50)  
);
```

### ► Insert

```
INSERT INTO Employees (id, name, salary, department)  
VALUES (1, 'Alice', 60000, 'HR');
```

### ► Select

```
SELECT * FROM Employees;  
SELECT name, salary FROM Employees;  
SELECT * FROM Employees WHERE salary > 50000;
```

### ► Update

```
UPDATE Employees SET salary = 70000 WHERE id = 1;
```

► **Delete**

```
DELETE FROM Employees WHERE id = 1;
```

---

#### 4. FILTERING DATA

► **WHERE, AND, OR, NOT**

```
SELECT * FROM Employees  
WHERE department = 'HR' AND salary > 50000;
```

► **IN, BETWEEN, LIKE**

```
SELECT * FROM Employees WHERE department IN ('HR', 'Sales');  
SELECT * FROM Employees WHERE salary BETWEEN 40000 AND 70000;  
SELECT * FROM Employees WHERE name LIKE 'A%'; -- starts with A
```

---

#### 5. SORTING & LIMITING

```
SELECT * FROM Employees ORDER BY salary DESC;  
SELECT * FROM Employees LIMIT 5;
```

---

#### 6. AGGREGATE FUNCTIONS

```
SELECT COUNT(*) FROM Employees;  
SELECT AVG(salary) FROM Employees;  
SELECT SUM(salary), MAX(salary), MIN(salary) FROM Employees;
```

---

#### 7. GROUPING & FILTERING AGGREGATES

```
SELECT department, AVG(salary) AS avg_salary  
FROM Employees  
GROUP BY department  
HAVING AVG(salary) > 50000;
```

---

#### 8. JOINS

► **Syntax:**

```
SELECT e.name, d.name  
FROM Employees e
```

JOIN Departments d ON e.department\_id = d.id;

### 8.1. INNER JOIN

```
SELECT e.name, d.dept_name
FROM Employees e
INNER JOIN Departments d ON e.department_id = d.id;
```

- ✓ Only matching department IDs shown.

### 8.2. LEFT JOIN

```
SELECT e.name, d.dept_name
FROM Employees e
LEFT JOIN Departments d ON e.department_id = d.id;
```

- ✓ All employees shown, even if no department.

### 8.3. RIGHT JOIN

```
SELECT e.name, d.dept_name
FROM Employees e
RIGHT JOIN Departments d ON e.department_id = d.id;
```

- ✓ All departments shown, even if no employee.

### 8.4. FULL OUTER JOIN

```
SELECT e.name, d.dept_name
FROM Employees e
LEFT JOIN Departments d ON e.department_id = d.id
UNION
SELECT e.name, d.dept_name
FROM Employees e
RIGHT JOIN Departments d ON e.department_id = d.id;
```

---

## 9. SUBQUERIES

Used to embed one query inside another.

```
SELECT name FROM Employees
WHERE salary > (
    SELECT AVG(salary) FROM Employees
);
```

---

## 10. CONSTRAINTS

- **PRIMARY KEY:** Uniquely identifies each record
  - **FOREIGN KEY:** Links to primary key of another table
  - **UNIQUE:** Ensures all values in a column are different
  - **NOT NULL:** Disallows NULL values
  - **CHECK:** Ensures a condition is met
- 

## 11. NORMALIZATION

- Process of organizing data to reduce redundancy and improve data integrity.
- **1NF:** Atomic values
- **2NF:** No partial dependencies
- **3NF:** No transitive dependencies
- Ensures data integrity and minimizes duplication.

### 🎯 Why Normalize?

- Eliminate duplicate data
  - Ensure data consistency
  - Improve data structure for querying and updates
- 

### 📁 Normal Forms (NF)

#### ✓ 1NF – First Normal Form

**Rule:** All values must be atomic (indivisible).

**Fix:** Remove repeating groups and store one value per cell.

#### Example (Bad):

**ID Name    Phones**

1 Alice    1234, 5678

#### Fix (1NF):

**ID Name Phone**

1 Alice    1234

1 Alice    5678

---

## ✓ 2NF – Second Normal Form

**Rule:** Be in 1NF + No Partial Dependency on a composite key.

**Fix:** Move partially dependent data to a new table.

**Example:**

Composite key: (StudentID, CourseID)

**StudentID CourseID StudentName**

**Fix (2NF):**

- Table 1: StudentCourses(StudentID, CourseID)
  - Table 2: Students(StudentID, StudentName)
- 

## ✓ 3NF – Third Normal Form

**Rule:** Be in 2NF + No transitive dependency.

**Fix:** Remove data that's indirectly dependent on the primary key.

**Example (Bad):**

**EmpID Name DeptID DeptName**

**Fix (3NF):**

- Table 1: Employees(EmpID, Name, DeptID)
  - Table 2: Departments(DeptID, DeptName)
- 

## □ When Not to Normalize?

- For read-heavy applications (analytics, reporting)
  - When query performance is more important than data integrity
- 

## ✓ Real-Life Analogy

- **1NF:** Each house (record) has its own mailbox (cell), not a shared one
- **2NF:** One key opens only one mailbox (no shared access)

- **3NF:** The key opens the right mailbox and not someone else's cabinet (no indirect dependencies)
- 

## 12. INDEXING

**Indexing** is a powerful technique to speed up `SELECT` queries on large tables.

- Speeds up `WHERE`, `JOIN`, `GROUP BY`, and `ORDER BY` queries
- Boosts search performance on large datasets.
- Avoid on frequently updated or low-cardinality columns
- Use `EXPLAIN` to test effectiveness
- Avoid over-indexing (inserts & updates become slower).

### ✓ Types

- **Single-column Index**

```
CREATE INDEX idx_salary ON Employees(salary);
```

- **Composite Index**

```
CREATE INDEX idx_name_dept ON Employees(name, department);
```

---

## 13. VIEWS

- A **view** is a virtual table based on a query
- They do not store data. Use **Materialized Views** (if supported) for cached result.
- Simplify complex queries.
- Enhance security (restrict column access)
- Enable reusable logic

### ✓ Example

```
CREATE VIEW HR_Employees AS
SELECT name, salary
FROM Employees
WHERE department = 'HR';
```

---

## 14. TRANSACTIONS & ACID

A **Transaction** is a logical unit of work consisting of one or more SQL statements.

Use `ROLLBACK;` to undo changes before `COMMIT;`.

- **ACID:**
  - **Atomicity:** All or none
  - **Consistency:** Valid data state
  - **Isolation:** No interference
  - **Durability:** Permanent changes

#### ✓ Example

```
BEGIN;  
UPDATE Accounts SET balance = balance - 100 WHERE id = 1;  
UPDATE Accounts SET balance = balance + 100 WHERE id = 2;  
COMMIT;
```

---

## 15. WINDOW FUNCTIONS

**Window functions** perform calculations across rows related to the current row — without collapsing them like `GROUP BY`.

Used for row-level calculations without grouping.

Use cases : Leaderboards, Change Tracking, Tired analysis

**RANK(), DENSE\_RANK(), ROW\_NUMBER()**

#### ✓ Example

```
SELECT name, salary,  
       RANK() OVER (ORDER BY salary DESC) AS salary_rank  
FROM Employees;
```

#### 🔑 Popular Window Functions

- `ROW_NUMBER()`, `RANK()`, `DENSE_RANK()`
  - `LAG()`, `LEAD()`
  - `NTILE(n)` (for percentiles)
- 

## 16. COMMON TABLE EXPRESSIONS (CTEs)

**CTEs** are temporary result sets used within queries — easier to read & maintain.

#### ✓ Syntax

```
WITH HighEarners AS (
```

```
SELECT * FROM Employees WHERE salary > 50000
)
SELECT name FROM HighEarnings;
```

#### □ Benefits

- Break complex queries into readable blocks
  - Use recursive CTEs for hierarchical data
- 

## 17. UNION vs UNION ALL

```
SELECT name FROM Managers
UNION
SELECT name FROM Employees; -- Removes duplicates
```

```
SELECT name FROM Managers
UNION ALL
SELECT name FROM Employees; -- Keeps duplicates
```

---

## 18. SET OPERATIONS

```
SELECT name FROM A
INTERSECT
SELECT name FROM B;
```

```
SELECT name FROM A
EXCEPT
SELECT name FROM B;
```

- ⚠ Not supported in MySQL; supported in PostgreSQL/SQL Server.
- 

## 19. NULL Handling

```
SELECT name, COALESCE(nickname, 'No Nick') FROM Users;
SELECT NULLIF(salary, 0); -- Returns NULL if salary = 0
```

- Use **IS NULL / IS NOT NULL**
  - **COALESCE(), IFNULL(), NULLIF()**
- 

## 20. CASE Statements

```
SELECT name,
```



```
CASE
  WHEN salary > 70000 THEN 'High'
  WHEN salary > 40000 THEN 'Medium'
  ELSE 'Low'
END AS salary_level
FROM Employees;
```

---

## 21. Stored Procedures & Functions (Advanced)

**Stored Procedures** are saved SQL blocks that perform actions.

**Functions** return a single value. Use procedures for reusable logic. Use functions in SELECT, WHERE, or JOIN clauses

### ✔ Procedure Example

```
CREATE PROCEDURE GetHighEarnings()
BEGIN
  SELECT * FROM Employees WHERE salary > 70000;
END;
```

### ✔ Function Example

```
CREATE FUNCTION TaxAmount(salary DECIMAL)
RETURNS DECIMAL
BEGIN
  RETURN salary * 0.10;
END;
```

---

## 22. Temporary Tables

```
CREATE TEMPORARY TABLE Temp_Employees AS
SELECT * FROM Employees WHERE department = 'HR';
```

---

## 23. Triggers (Advanced - Rare)

```
CREATE TRIGGER before_insert_trigger
BEFORE INSERT ON Employees
FOR EACH ROW
SET NEW.salary = IF(NEW.salary < 0, 0, NEW.salary);
```

- Used for validation, logging, auto-calculation.
- 

## 24. Data Definition vs Data Manipulation vs Data Control

- **DDL (Definition):** CREATE, ALTER, DROP
  - **DML (Manipulation):** SELECT, INSERT, UPDATE, DELETE
  - **DCL (Control):** GRANT, REVOKE
  - **TCL (Transaction):** COMMIT, ROLLBACK
- 

## INTERVIEW QUERY PATTERNS

### ► Nth Highest Salary

```
SELECT DISTINCT salary
FROM Employees
ORDER BY salary DESC
LIMIT 1 OFFSET 1; -- 2nd highest
```

### ► Duplicate Rows

```
SELECT name, COUNT(*)
FROM Employees
GROUP BY name
HAVING COUNT(*) > 1;
```

### ► Same Salary Employees

```
SELECT * FROM Employees
WHERE salary IN (
    SELECT salary FROM Employees
    GROUP BY salary
    HAVING COUNT(*) > 1
);
```

### ► Employees with Max Salary Per Department

```
SELECT name, department_id, salary
FROM (
    SELECT *,
        RANK() OVER (PARTITION BY department_id ORDER BY salary DESC) AS rnk
    FROM Employees
) ranked
WHERE rnk = 1;
```

### ► Self Join Example

```
SELECT A.name AS Employee, B.name AS Manager
FROM Employees A
JOIN Employees B ON A.manager_id = B.id;
```

## ► EXISTS vs IN vs JOIN

```
-- EXISTS
SELECT name FROM Employees e
WHERE EXISTS (
  SELECT 1 FROM Departments d WHERE d.id = e.department_id
);
-- IN
SELECT name FROM Employees
WHERE department_id IN (SELECT id FROM Departments);
-- JOIN
SELECT e.name, d.name FROM Employees e
JOIN Departments d ON e.department_id = d.id;
```

---

# Advanced SQL Topics

## 1. Recursive CTEs (Made Simple)

A **Recursive CTE** lets a query call itself, useful when working with **hierarchies** like employees & managers or categories & subcategories.

### Easy Analogy:

Think of it like a loop in SQL. Start with one person (manager), then find their subordinates, then subordinates of subordinates, and so on.

### Syntax:

```
WITH RECURSIVE cte_name AS (  
  SELECT ... -- Starting point (anchor)  
  UNION ALL  
  SELECT ... FROM cte_name ... -- Repeat until done  
)  
SELECT * FROM cte_name;
```

### Example:

```
WITH RECURSIVE Subordinates AS (  
  SELECT id, name, manager_id FROM Employees WHERE id = 1 -- Top manager  
  UNION ALL  
  SELECT e.id, e.name, e.manager_id  
  FROM Employees e  
  JOIN Subordinates s ON e.manager_id = s.id  
)  
SELECT * FROM Subordinates;
```

✓ This will return all employees working under manager ID 1, even indirectly.

---

## 2. Window Functions

These functions **look at other rows** in the result **without grouping** them. Useful when you want to compare each row with others.

### Everyday Use Case:

"Show each employee's salary along with the previous and next employee's salary."

### Key Functions:

- LAG(column) – Looks **before** current row
- LEAD(column) – Looks **after** current row
- NTILE(n) – Divides into n equal groups (for percentiles/quartiles)
- RANK(), DENSE\_RANK(), ROW\_NUMBER() – Ranking logic

#### Example:

```
SELECT name, salary,  
       LAG(salary) OVER (ORDER BY salary) AS prev_salary,  
       LEAD(salary) OVER (ORDER BY salary) AS next_salary  
FROM Employees;
```

---

### 3. Working with JSON in SQL

Modern databases allow storing and querying **JSON data** inside SQL columns (especially useful in dynamic or semi-structured data).

#### PostgreSQL:

```
SELECT data->>'name' AS name FROM orders WHERE data->>'status' = 'shipped';
```

#### MySQL:

```
SELECT JSON_EXTRACT(data, '$.name') AS name FROM orders;
```

🔍 Here, data is a column containing JSON like: { "name": "Laptop", "status": "shipped" }

---

### 4. Performance Optimization Techniques

Optimizing your SQL queries can **make your app faster**, especially for large data.

#### a. Indexing

- Like a book index: helps locate rows faster
- Useful for WHERE, JOIN, and ORDER BY

```
CREATE INDEX idx_salary ON Employees(salary);
```

#### b. Query Refactoring

- Avoid SELECT \* → only select needed columns
- Use EXISTS instead of IN when subquery is large
- Apply filters early (use WHERE before GROUP BY)

#### c. Execution Plan

Use EXPLAIN to understand how the database runs your query:

```
EXPLAIN SELECT * FROM Employees WHERE department_id = 3;
```

---

## 5. Stored Procedures & Functions (Simplified)

Stored Procedures = SQL scripts stored in the database that can be reused. Like a saved function.

### Example with Loop:

```
DELIMITER //
CREATE PROCEDURE GiveBonus()
BEGIN
  DECLARE done INT DEFAULT FALSE;
  DECLARE emp_id INT;
  DECLARE emp_cursor CURSOR FOR SELECT id FROM Employees;
  DECLARE CONTINUE HANDLER FOR NOT FOUND SET done = TRUE;

  OPEN emp_cursor;
  read_loop: LOOP
    FETCH emp_cursor INTO emp_id;
    IF done THEN
      LEAVE read_loop;
    END IF;
    UPDATE Employees SET salary = salary + 5000 WHERE id = emp_id;
  END LOOP;
  CLOSE emp_cursor;
END //
DELIMITER ;
```

✓ This procedure gives a ₹5000 bonus to all employees.

---

## 6. Importing/Exporting Data

Moving data between CSV files and SQL is common in real projects.

### MySQL:

```
LOAD DATA INFILE '/path/file.csv'
INTO TABLE Employees
FIELDS TERMINATED BY ','
LINES TERMINATED BY '\n'
IGNORE 1 ROWS;
```

### PostgreSQL:

```
COPY Employees(name, salary)
FROM '/path/file.csv'
DELIMITER ','
CSV HEADER;
```

---

## 7. Role-Based Access Control (Security Basics)

SQL lets you give different permissions to different users.

### Example:

```
GRANT SELECT, INSERT ON Employees TO 'analyst';
REVOKE DELETE ON Employees FROM 'analyst';
```

✓ 'analyst' can view and add data, but can't delete anything.

---

## 8. Star vs Snowflake Schema

Used in **Data Warehouses & Reporting Tools**.

### Star Schema:

- One central fact table (e.g., Sales)
- Linked to dimension tables (Customer, Product)
- Easy to query, faster for reporting

### Snowflake Schema:

- Dimensions are normalized
  - Less redundancy but more joins
- 

## 9. OLTP vs OLAP (Database Types Simplified)

| Feature  | OLTP                     | OLAP                      |
|----------|--------------------------|---------------------------|
| Purpose  | Daily transactions       | Data analysis             |
| Design   | Highly normalized tables | Denormalized schema       |
| Examples | Banking, eCommerce       | Reporting, Data Warehouse |

---

## INTERVIEW & PLACEMENT

### ◆ Basic Level (Entry / Fresher)

1. **What is SQL?**  
Structured Query Language used to interact with relational databases.
2. **What is the difference between WHERE and HAVING?**  
WHERE filters rows before aggregation, HAVING filters after aggregation.
3. **What is a Primary Key?**  
A unique identifier for each row in a table. Cannot contain NULL values.
4. **Difference between INNER JOIN and LEFT JOIN?**  
INNER JOIN: Returns matched rows only.  
LEFT JOIN: Returns all rows from the left table, even if no match.
5. **How do you handle NULLs in SQL?**  
Use IS NULL, IS NOT NULL, COALESCE(), IFNULL(), NULLIF().
6. **What is a Foreign Key?**  
It links two tables by referring to the primary key of another table.
7. **What is normalization?**  
Process of organizing data to reduce redundancy and improve integrity.
8. **What is denormalization?**  
Combining tables to improve read performance by reducing joins.
9. **What does the DISTINCT keyword do?**  
Removes duplicate rows from the result set.
10. **What is the use of LIMIT or TOP?**  
Used to limit the number of rows returned by a query.

### ◆ Intermediate Level

11. **What are aggregate functions?**  
Functions like SUM(), AVG(), COUNT(), MAX(), MIN().
12. **What is a subquery?**  
A query nested inside another query.
13. **What is the difference between UNION and UNION ALL?**  
UNION removes duplicates, UNION ALL includes all rows.
14. **Explain CASE statement.**  
Used to apply conditional logic within queries.
15. **What are indexes in SQL?**  
They improve the speed of data retrieval.
16. **What is a view?**  
A virtual table based on the result-set of a query.
17. **What are window functions?**  
Functions like RANK(), ROW\_NUMBER() that work across rows.
18. **Difference between RANK() and DENSE\_RANK()?**  
RANK() skips ranks on ties, DENSE\_RANK() does not.
19. **What is a CTE (Common Table Expression)?**  
A temporary result set defined within the execution scope of a query.
20. **Explain ACID properties.**  
Atomicity, Consistency, Isolation, Durability - ensures reliable transactions.



21. **What is a composite key?**  
A primary key made of multiple columns.
22. **Difference between DELETE and TRUNCATE?**  
DELETE can be conditional and logs row-by-row deletion; TRUNCATE removes all rows faster without logging each deletion.
23. **What is a surrogate key?**  
A unique identifier for an entity that is not derived from application data.
24. **What is referential integrity?**  
Ensures foreign key values match primary key values in the referenced table.
25. **How does EXISTS differ from IN?**  
EXISTS stops on first match; IN evaluates all results.
26. **How do you use GROUP BY with multiple columns?**  
You can group by multiple columns by separating them with commas.
27. **Can we use ORDER BY with GROUP BY?**  
Yes. GROUP BY groups the data; ORDER BY sorts the grouped results.
28. **What is a scalar subquery?**  
A subquery that returns exactly one value.
29. **What are correlated subqueries?**  
Subqueries that refer to columns from the outer query.
30. **What is the use of ISNULL() or IFNULL()?**  
To replace NULL values with custom values.

#### ◆ Advanced Level

31. **What are triggers in SQL?**  
Procedures that automatically execute on certain events.
32. **What are stored procedures?**  
Reusable blocks of SQL statements stored in the database.
33. **What is the use of EXPLAIN or EXPLAIN PLAN?**  
To understand how the database executes a query.
34. **What are transactions?**  
A unit of work that is performed against a database.
35. **How do you optimize a slow query?**  
Use indexing, limit joins, avoid SELECT \*, and use EXPLAIN to analyze.
36. **Difference between clustered and non-clustered index?**  
Clustered index determines row order in the table; non-clustered does not.
37. **Difference between OLTP and OLAP systems?**  
OLTP: Online Transaction Processing (day-to-day operations).  
OLAP: Online Analytical Processing (data analysis and reporting).
38. **What are materialized views?**  
Stored query results that can be refreshed periodically.
39. **How do you implement pagination in SQL?**  
Using LIMIT and OFFSET or ROW\_NUMBER() for custom logic.
40. **How do you handle duplicate rows?**  
Using ROW\_NUMBER() or DISTINCT or CTEs with filtering.
41. **What are the different types of joins?**  
INNER, LEFT, RIGHT, FULL OUTER, CROSS JOIN, SELF JOIN.
42. **What is a CROSS JOIN?**  
Returns Cartesian product of two tables.

**43. Can you sort by an alias in SQL?**

Yes, you can use the alias name in the ORDER BY clause.

**44. What is the difference between SQL and NoSQL?**

SQL is relational and uses tables. NoSQL is non-relational and uses documents, key-value pairs, etc.

**45. What are the common data types in SQL?**

INT, VARCHAR, DATE, BOOLEAN, DECIMAL, TEXT.

**46. Can a table have multiple foreign keys?**

Yes, a table can reference multiple other tables using foreign keys.

**47. What happens if you violate a foreign key constraint?**

The query fails with an integrity constraint violation.

**48. How can you change a column datatype in SQL?**

Using ALTER TABLE table\_name MODIFY column\_name new\_datatype;

**49. What are NULL-safe operators?**

Operators like <=> in MySQL allow safe comparison with NULL.

**50. How would you detect and remove duplicate records?**

Using CTE and ROW\_NUMBER() to filter duplicates.

---

# Real Company SQL Questions

## Google

- Second Highest Salary

```
SELECT MAX(salary) AS SecondHighest
FROM Employees
WHERE salary < (SELECT MAX(salary) FROM Employees);
```

---

- Remove Duplicate Emails Using ROW\_NUMBER()

```
WITH RankedEmails AS (
  SELECT *, ROW_NUMBER() OVER (PARTITION BY email ORDER BY id)
  AS rn
  FROM Users
)
DELETE FROM RankedEmails WHERE rn > 1;
```

---

- Employees Earning More Than Average Salary

```
SELECT name, salary
FROM Employees
WHERE salary > (SELECT AVG(salary) FROM Employees);
```

---

## Amazon

- Customers Who Ordered in Every Month

```
SELECT customer_id
FROM Orders
GROUP BY customer_id
HAVING COUNT(DISTINCT MONTH(order_date)) = 12;
```

---

- Top 3 Earners in Each Department

```
SELECT *
FROM (
  SELECT *, DENSE_RANK() OVER (PARTITION BY department_id
  ORDER BY salary DESC) AS rank
  FROM Employees
) ranked
WHERE rank <= 3;
```

---

- Find Returning Users Using LAG

```
SELECT user_id, order_date,  
       LAG(order_date) OVER (PARTITION BY user_id ORDER BY  
order_date) AS previous_order  
FROM Orders;
```

---

## TCS

- Employees Joined in Last 6 Months

```
SELECT * FROM Employees  
WHERE join_date >= DATE_SUB(CURDATE(), INTERVAL 6 MONTH);
```

---

- Manager-Reportee Mapping

```
SELECT e.name AS Employee, m.name AS Manager  
FROM Employees e  
JOIN Employees m ON e.manager_id = m.id;
```

---

- Department with Highest Average Salary

```
SELECT department_id  
FROM Employees  
GROUP BY department_id  
ORDER BY AVG(salary) DESC  
LIMIT 1;
```

---

## Flipkart

- Top 5 Selling Products by Revenue

```
SELECT product_id, SUM(price * quantity) AS revenue  
FROM Orders  
GROUP BY product_id  
ORDER BY revenue DESC  
LIMIT 5;
```

---

- Users with More Than 3 Failed Transactions

```
SELECT user_id  
FROM Transactions  
WHERE status = 'FAILED'
```

```
GROUP BY user_id  
HAVING COUNT(*) > 3;
```

---

- Orders Not Yet Shipped

```
SELECT *  
FROM Orders  
WHERE status = 'PLACED' AND shipped_date IS NULL;
```

---

### SQL Problem-Solving Round

- Nth Highest Salary

```
SELECT DISTINCT salary  
FROM Employees  
ORDER BY salary DESC  
LIMIT 1 OFFSET N-1;
```

---

- Duplicate Emails

```
SELECT email, COUNT(*)  
FROM Users  
GROUP BY email  
HAVING COUNT(*) > 1;
```

---

- Top Earner per Department

```
SELECT name, department_id, salary  
FROM (  
    SELECT *, RANK() OVER (PARTITION BY department_id ORDER BY  
salary DESC) AS rnk  
    FROM Employees  
    ) ranked  
WHERE rnk = 1;
```

---

- Customers Who Never Ordered

```
SELECT c.customer_id, c.name  
FROM Customers c  
LEFT JOIN Orders o ON c.customer_id = o.customer_id  
WHERE o.order_id IS NULL;
```

---

- Employees Hired in the Last Month

```
SELECT * FROM Employees  
WHERE hire_date BETWEEN DATE_SUB(CURDATE(), INTERVAL 1  
MONTH) AND CURDATE();
```

---

- [User Order Summary](#)

```
SELECT user_id, COUNT(*) AS order_count, SUM(total_amount) AS  
total_spent  
FROM Orders  
GROUP BY user_id;
```

---

- [7-Day Rolling Login Count](#)

```
SELECT user_id, login_date,  
COUNT(*) OVER (  
    PARTITION BY user_id ORDER BY login_date  
    ROWS BETWEEN 6 PRECEDING AND CURRENT ROW  
) AS seven_day_logins  
FROM Logins;
```

---

- [Funnel Drop-off Analysis](#)

```
SELECT stage, COUNT(DISTINCT user_id) AS user_count  
FROM Funnel  
GROUP BY stage  
ORDER BY stage;
```

---

- [Average Time Between Orders](#)

```
SELECT user_id, AVG(DATEDIFF(order_date, LAG(order_date) OVER  
(PARTITION BY user_id ORDER BY order_date))) AS  
avg_days_between_orders  
FROM Orders;
```

---

- [Self Join for Employee Manager Mapping](#)

```
SELECT e.name AS Employee, m.name AS Manager  
FROM Employees e  
JOIN Employees m ON e.manager_id = m.id;
```

---